

## CHAPTER 14



# Transactions

### Practice Exercises

- 14.1 **Answer:** Even in this case the recovery manager is needed to perform roll-back of aborted transactions.
- 14.2 **Answer:** There are several steps in the creation of a file. A storage area is assigned to the file in the file system, a unique i-number is given to the file and an i-node entry is inserted into the i-list. Deletion of file involves exactly opposite steps.  
For the file system user in UNIX, durability is important for obvious reasons, but atomicity is not relevant generally as the file system doesn't support transactions. To the file system implementor though, many of the internal file system actions need to have transaction semantics. All the steps involved in creation/deletion of the file must be atomic, otherwise there will be unreferenceable files or unusable areas in the file system.
- 14.3 **Answer:** Database systems usually perform crucial tasks whose effects need to be atomic and durable, and whose outcome affects the real world in a permanent manner. Examples of such tasks are monetary transactions, seat bookings etc. Hence the ACID properties have to be ensured. In contrast, most users of file systems would not be willing to pay the price (monetary, disk space, time) of supporting ACID properties.
- 14.4 **Answer:** If a transaction is very long or when it fetches data from a slow disk, it takes a long time to complete. In absence of concurrency, other transactions will have to wait for longer period of time. Average response time will increase. Also when the transaction is reading data from disk, CPU is idle. So resources are not properly utilized. Hence concurrent execution becomes important in this case. However, when the transactions are short or the data is available in memory, these problems do not occur.
- 14.5 **Answer:** Most of the concurrency control protocols (protocols for ensuring that only serializable schedules are generated) used in practice are based on conflict serializability—they actually permit only a subset of

conflict serializable schedules. The general form of view serializability is very expensive to test, and only a very restricted form of it is used for concurrency control.

**14.6 Answer:** There is a serializable schedule corresponding to the precedence graph below, since the graph is acyclic. A possible schedule is obtained by doing a topological sort, that is,  $T_1, T_2, T_3, T_4, T_5$ .

**14.7 Answer:** A cascadeless schedule is one where, for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads data items previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ . Cascadeless schedules are desirable because the failure of a transaction does not lead to the aborting of any other transaction. Of course this comes at the cost of less concurrency. If failures occur rarely, so that we can pay the price of cascading aborts for the increased concurrency, noncascadeless schedules might be desirable.

**14.8 Answer:**

a. A schedule showing the Lost Update Anomaly:

$T_1$	$T_2$
<b>read(A)</b>	
	<b>read(A)</b>
<b>write(A)</b>	<b>write(A)</b>

In the above schedule, the value written by the transaction  $T_2$  is lost because of the write of the transaction  $T_1$ .

b. Lost Update Anomaly in Read Committed Isolation Level

$T_1$	$T_2$
<b>lock-S(A)</b>	
<b>read(A)</b>	
<b>unlock(A)</b>	
	<b>lock-X(A)</b>
	<b>read(A)</b>
	<b>write(A)</b>
	<b>unlock(A)</b>
	<b>commit</b>
<b>lock-X(A)</b>	
<b>write(A)</b>	
<b>unlock(A)</b>	
<b>commit</b>	

The locking in the above schedule ensures the Read Committed isolation level. The value written by transaction  $T_2$  is lost due to  $T_1$ 's write.

c. Lost Update Anomaly is not possible in Repeatable Read isolation level. In repeatable read isolation level, a transaction  $T_1$  reading a

data item  $X$ , holds a shared lock on  $X$  till the end. This makes it impossible for a newer transaction  $T_2$  to write the value of  $X$  (which requires  $X$ -lock) until  $T_1$  finishes. This forces the serialization order  $T_1, T_2$  and thus the value written by  $T_2$  is not lost.

**14.9 Answer:**

Suppose that the bank enforces the integrity constraint that the sum of the balances in the checking and the savings account of a customer must not be negative. Suppose the checking and savings balances for a customer are \$100 and \$200 respectively.

Suppose that transaction  $T_1$  withdraws \$200 from the checking account after verifying the integrity constraint by reading both the balances. Suppose that concurrent transaction  $T_2$  withdraws \$200 from the checking account after verifying the integrity constraint by reading both the balances.

Since each of the transactions checks the integrity constraints on its own snapshot, if they run concurrently each will believe that the sum of the balances after the withdrawal is \$100 and therefore its withdrawal does not violate the integrity constraint. Since the two transactions update different data items, they do not have any update conflict, and under snapshot isolation both of them can commit. This is a non-serializable execution which results into a serious problem of withdrawal of more money.

**14.10 Answer:**

Consider a web-based airline reservation system. There could be many concurrent requests to see the list of available flights and available seats in each flight and to book tickets. Suppose, there are two users  $A$  and  $B$  concurrently accessing this web application, and only one seat is left on a flight.

Suppose that both user  $A$  and user  $B$  execute transactions to book a seat on the flight, and suppose that each transaction checks the total number of seats booked on the flight, and inserts a new booking record if there are enough seats left. Let  $T_3$  and  $T_4$  be their respective booking transactions, which run concurrently. Now  $T_3$  and  $T_4$  will see from their snapshots that one ticket is available and insert new booking records. Since the two transactions do not update any common data item (tuple), snapshot isolation allows both transactions to commit. This results in an extra booking, beyond the number of seats available on the flight.

However, this situation is usually not very serious since cancellations often resolve the conflict; even if the conflict is present at the time the flight is to leave, the airline can arrange a different flight for one of the passengers on the flight, giving incentives to accept the change. Using snapshot isolation improves the overall performance in this case since the booking transactions read the data from their snapshots only and do not block other concurrent transactions.

## 14.11 Answer:

The given situation will not cause any problem for the definition of conflict serializability since the ordering of operations on each data item is necessary for conflict serializability, whereas the ordering of operations on different data items is not important.

$T_1$	$T_2$
<b>read</b> (A)	
<b>write</b> (B)	<b>read</b> (B)

For the above schedule to be conflict serializable, the only ordering requirement is **read**(B)  $\rightarrow$  **write**(B). **read**(A) and **read**(B) can be in any order. Therefore, as long as the operations on a data item can be totally ordered, the definition of conflict serializability should hold on the given multi-processor system.